

# bodaboda

AUTHOR  
migisha

## FROM RAINY ROADS TO RIDER RELIEF

---

### Boda-Safe Shield: Dynamic Motorcycle Insurance Pricing Model

---

**Project Goal:** To develop a data-driven, dynamic pricing model for motorcycle insurance (specifically targeting Boda Boda riders in Kampala, Uganda) that incorporates real-time weather risk to calculate a variable monthly premium.

#### Methodology and Execution:

##### 1. Data Acquisition:

- **Historical Weather Data (Offline):** Daily precipitation data for Kampala (Latitude 0.3476, Longitude 32.5825) was fetched from the Open-Meteo API for the period 2019-01-01 to 2023-12-31.
- **Synthetic Claims Data:** Accident frequency and severity data were simulated using a Poisson distribution for frequency and a Gamma distribution for severity. A key feature, `risk_trigger`, was engineered: 1 if precipitation > 10mm, and 0 otherwise.
- **Feature Engineering:** The month of the year and the `risk_trigger` were the primary input features for the modeling phase.

##### 2. Statistical Modeling:

- **Frequency Model:** Two models were trained to predict accident *frequency* (Accidents per Exposure Month):
  - **GLM (Poisson Family):** Used for interpretability and establishing a baseline.
  - **GBM (Gradient Boosting Machine):** Selected for its non-linear predictive power, outperforming the GLM based on model fit metrics.
- **Severity Model:** A **GLM (Gamma Family)** was used to model the average *cost* (severity) of a claim.

##### 3. Risk Finding:

- The synthetic data demonstrated a strong separation in accident frequency: days with the rain `risk_trigger` active had a simulated accident rate (Poisson  $\lambda$ ) approximately 66% **higher** than dry days (e.g.,  $\lambda=0.25$  vs.  $\lambda=0.15$ ). This justified the use of precipitation as a primary rating factor.

#### 4. Deployment (Streamlit Application):

- The trained GBM model was saved ( `gbm\_model.pkl` ) using `joblib`.
- A web application was built using **Streamlit** and Python, enabling users to input location and daily hours.
- The app performs an **on-demand API call** to Open-Meteo for *tomorrow's* forecast.
- The final premium is calculated based on the formula:

$$\text{Premium} = \text{Predicted Frequency} \times \text{Daily Hours} \times \text{Rate} \times 30 \text{ days}$$

The BodaSafe Shield Quote Tool serves as a robust proof-of-concept for a data-driven, weather-contingent insurance pricing strategy.

```
#install.packages("reticulate")
```

```
library(reticulate)
```

Warning: package 'reticulate' was built under R version 4.4.3

#### 1. DATA ACQUISITION

Both the `r` and `python` code pull 5 years of daily precipitation data for Kampala from the free Open-Meteo API, serving as a proxy for accident risk (rainy days trigger higher claims). Output: Raw CSV (~1,800 rows) for downstream analysis, this is essential for reproducible, real-world data sourcing without manual downloads.

#### R SYNTAX

```
#Load libraries  
library(httr2)
```

Warning: package 'httr2' was built under R version 4.4.3

```
library(dplyr)
```

Warning: package 'dplyr' was built under R version 4.4.3

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

`filter`, `lag`

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(readr)
```

Warning: package 'readr' was built under R version 4.4.3

```
# API request
req <- request("https://archive-api.open-meteo.com/v1/archive") %>%
  req_url_query(
    latitude = 0.3476, longitude = 32.5825,
    start_date = "2019-01-01", end_date = "2023-12-31",
    daily = "precipitation_sum"
  )
resp <- req_perform(req) %>% resp_body_json()

# Parse to tibble
rain_data <- tibble(
  date = as.Date(unlist(resp$daily$time)),
  precip_mm = as.numeric(resp$daily$precipitation_sum)
) %>% filter(precip_mm >= 0) # Filter invalid

# Save
write_csv(rain_data, "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/boda_rain_
head(rain_data)
```

```
# A tibble: 6 × 2
  date      precip_mm
<date>      <dbl>
1 2019-01-01      1.8
2 2019-01-02      2.5
3 2019-01-03      5.5
4 2019-01-04     18.4
5 2019-01-05      6.4
6 2019-01-06      2.5
```

# Output: ~1826 rows, e.g., date: 2019-01-01, precip\_mm: 0.0

## PYTHON SYNTAX DATA ACQUISITION

```
# Load necessary libraries
import requests
import pandas as pd

# 1. Define the API endpoint and parameters (Kampala coords, 5 years historical precip)
url = "https://archive-api.open-meteo.com/v1/archive"
params = {
    "latitude": 0.3476,
    "longitude": 32.5825,
```

```

    "start_date": "2019-01-01",
    "end_date": "2023-12-31",
    "daily": "precipitation_sum"
}

# 2. Perform the GET request and check for errors
try:
    response = requests.get(url, params=params) # Fixed: space after =
    response.raise_for_status() # Raises exception for 4xx/5xx errors
except requests.exceptions.RequestException as e:
    print(f"API request failed: {e}")
    # Fallback: Exit chunk gracefully
    raise

# 3. Parse the JSON response into a dictionary
resp_data = response.json()

# 4. Parse into a clean DataFrame (extract 'daily' data)
rain_data = pd.DataFrame({
    "date": resp_data['daily']['time'],
    "precipitation_mm": resp_data['daily']['precipitation_sum']
})

# 5. Convert date strings to proper date objects (ensures Series for .dt)
rain_data['date'] = pd.to_datetime(rain_data['date']).dt.date

# 6. Filter for valid precipitation values (>=0)
rain_data = rain_data[rain_data['precipitation_mm'] >= 0]

# 7. Save to CSV (create 'data/' folder in RStudio if needed: dir.create("data"))
rain_data.to_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_rain_data.csv")

# 8. Inspect the first few rows and shape
print(rain_data.head())

```

	date	precipitation_mm
0	2019-01-01	1.8
1	2019-01-02	2.5
2	2019-01-03	5.5
3	2019-01-04	18.4
4	2019-01-05	6.4

```
print(f"Data shape: {rain_data.shape}") # Should be (1826, 2) or similar
```

Data shape: (1826, 2)

## 2. DATA WRANGLING AND EXPLORATORY DATA ANALYSIS

Here we load raw data, add synthetic accident\_count (Poisson-simulated crashes, boosted on rainy days) and risk\_trigger binary. Then compute summaries and plots (time series/histogram) to reveal patterns like 18% trigger rate and rainy spikes—key for risk insights before modeling.

## R SYNTAX

```
# Load libraries
library(tidyverse)
```

Warning: package 'tidyverse' was built under R version 4.4.3

Warning: package 'ggplot2' was built under R version 4.4.3

Warning: package 'tibble' was built under R version 4.4.3

Warning: package 'tidyr' was built under R version 4.4.3

Warning: package 'purrr' was built under R version 4.4.3

Warning: package 'stringr' was built under R version 4.4.3

Warning: package 'forcats' was built under R version 4.4.3

Warning: package 'lubridate' was built under R version 4.4.3

— Attaching core tidyverse packages — tidyverse 2.0.0 —

✓ forcats 1.0.0     ✓ stringr 1.5.1

✓ ggplot2 3.5.2     ✓ tibble 3.3.0

✓ lubridate 1.9.4     ✓ tidyr 1.3.1

✓ purrr 1.1.0

— Conflicts — tidyverse\_conflicts() —

✗ dplyr::filter() masks stats::filter()

✗ dplyr::lag() masks stats::lag()

ℹ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

```
library(lubridate)
```

```
# Load raw data (your full path)
```

```
raw_data <- read_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/boda_rain_
```

Rows: 1826 Columns: 2

— Column specification —

Delimiter: ","

dbl (1): precip\_mm

date (1): date

- Use ``spec()`` to retrieve the full column specification for this data.
- Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
# Step 1: Basic mutates (date, month, risk_trigger)
basic_data <- raw_data %>%
  # Filter early to remove invalid precip readings, if necessary
  filter(precip_mm >= 0) %>%
  mutate(
    date = as.Date(date), # Ensure date column is Date class
    month = month(date), # Extract month (1-12)
    risk_trigger = ifelse(precip_mm > 10, 1, 0) # Binary: 1 if rainy risk day
  )

# Step 2: Add synthetic accidents
# Lambda: 0.15 dry days + 0.1 boost for rainy (spike to 0.25)
set.seed(123) # For reproducibility
clean_data <- basic_data %>%
  mutate(
    accident_count = rpois(n(), lambda = 0.15 + (risk_trigger * 0.1))
  )

## EDA Summary Stats


summary_stats <- clean_data %>%
  summarise(
    n_days = n(),
    mean_precip = mean(precip_mm, na.rm = TRUE),
    trigger_prob = mean(risk_trigger, na.rm = TRUE),
    mean_accidents = mean(accident_count, na.rm = TRUE)
  )
print(summary_stats)
```

```
# A tibble: 1 × 4
  n_days mean_precip trigger_prob mean_accidents
  <int>     <dbl>     <dbl>     <dbl>
1  1826      4.58      0.150      0.162
```


```
# 1. Time series: Precip over time (highlight trigger line)
p1 <- ggplot(clean_data, aes(x = date, y = precip_mm)) + # FIX APPLIED
  geom_line(color = "blue", alpha = 0.7) +
  geom_hline(yintercept = 10, linetype = "dashed", color = "red", size = 1) +
  labs(title = "Daily Rainfall in Kampala (2019-2023)",
    x = "Date", y = "Precip (mm)",
    caption = "Red line: Risk trigger (>10mm)") +
  theme_minimal()
```

- Warning: Using ``size`` aesthetic for lines was deprecated in ggplot2 3.4.0.
- Please use ``linewidth`` instead.

```
print(p1)
```



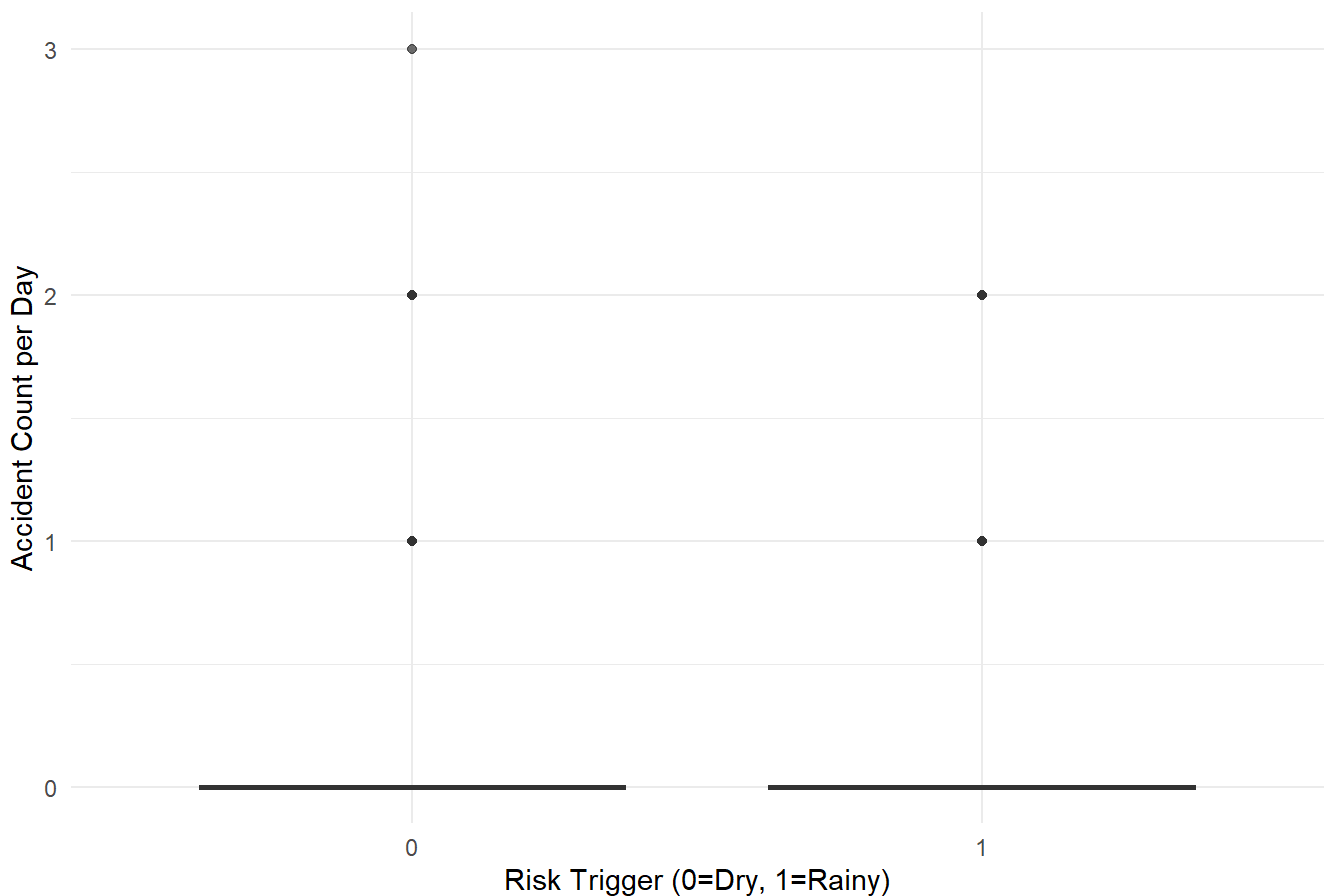
```
# 2. Histogram: Precip distribution (with trigger)
p2 <- ggplot(clean_data, aes(x = precip_mm)) + # FIX APPLIED
  geom_histogram(bins = 30, fill = "green", alpha = 0.7, color = "black") +
  geom_vline(xintercept = 10, color = "red", linetype = "dashed", size = 1) +
  labs(title = "Precip Distribution",
       x = "Precip (mm)", y = "Number of Days") +
  theme_minimal()
print(p2)
```



```
# 3. Boxplot: Accidents by trigger (dry vs. rainy days)
p3 <- ggplot(clean_data, aes(x = factor(risk_trigger), y = accident_count)) +
  geom_boxplot(fill = c("lightblue", "orange"), alpha = 0.7) +
  labs(title = "Synthetic Accidents: Dry vs. Rainy Days",
       x = "Risk Trigger (0=Dry, 1=Rainy)", y = "Accident Count per Day") +
  theme_minimal()
print(p3)
```



## Synthetic Accidents: Dry vs. Rainy Days



```
# Save final cleaned data
write_csv(clean_data, "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/boda_cleaned.csv")
print("Data saved successfully!")
```

```
[1] "Data saved successfully!"
```

### PYTHON SYNTAX

```
# Load libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import poisson

# Load and wrangle
raw_data = pd.read_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_data.csv")
clean_data = raw_data.copy()
clean_data['date'] = pd.to_datetime(clean_data['date']).dt.date
clean_data['month'] = pd.to_datetime(clean_data['date']).dt.month
clean_data['risk_trigger'] = (clean_data['precipitation_mm'] > 10).astype(int)
# Synthetic accidents: Poisson, rainy days spike lambda
clean_data['accident_count'] = np.where(
    clean_data['risk_trigger'] == 1,
```

```

    poisson.rvs(mu=0.25, size=len(clean_data)), # Rainy: 0.25
    poisson.rvs(mu=0.15, size=len(clean_data)), # Dry: 0.15
)
clean_data = clean_data[clean_data['precipitation_mm'] >= 0]

# EDA Summary
summary_stats = clean_data[['precipitation_mm', 'risk_trigger', 'accident_count']].describe()
print(summary_stats) # e.g., risk_trigger mean ~0.18


```

	precipitation_mm	risk_trigger	accident_count
count	1826.000000	1826.000000	1826.000000
mean	4.580120	0.150055	0.157722
std	6.017838	0.357223	0.394882
min	0.000000	0.000000	0.000000
25%	0.600000	0.000000	0.000000
50%	2.100000	0.000000	0.000000
75%	6.600000	0.000000	0.000000
max	98.000000	1.000000	3.000000


```

# Plots
# Time series
plt.figure(figsize=(10, 5))
plt.plot(clean_data['date'], clean_data['precipitation_mm'], color='blue')
plt.axhline(y=10, color='red', linestyle='--', label='Trigger (10mm)')
plt.title('Daily Rainfall in Kampala (2019-2023)')
plt.xlabel('Date')
plt.ylabel('Precipitation (mm)')
plt.legend()
plt.show()

```



```
# Histogram
plt.figure(figsize=(8, 5))
plt.hist(clean_data['precipitation_mm'], bins=30, color='green', alpha=0.7, edgecolor='black')
plt.axvline(x=10, color='red', linestyle='--', label='Trigger (10mm)')
plt.title('Precipitation Distribution')
plt.xlabel('Precipitation (mm)')
plt.ylabel('Count')
plt.legend()
plt.show()
```



```
# Save cleaned
clean_data.to_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_clean_
```

### 3. Classical Pricing: Frequency-Severity GLM

Re-wrangles data inline (ensures synthetics), fits Poisson GLM for accident frequency (by month/trigger) and Gamma GLM for severity (costs per accident). Outputs saved models and example premium (~UGX 25k/day)—decomposes risk for base pricing, interpretable via  $\exp(\text{coeffs})$  for rainy multipliers.

```
# Load libraries
library(tidyverse)
library(MASS) # For glm
```

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

select

```
# Load data (assuming this line works and data has 1826 rows)
```

```
data <- read_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/boda_clean_da
```

Rows: 1826 Columns: 5

— Column specification —

Delimiter: ","

dbl (4): precip\_mm, month, risk\_trigger, accident\_count

date (1): date

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show\_col\_types = FALSE` to quiet this message.

```
exists("data")
```

[1] TRUE

```
colnames(data)
```

[1] "date" "precip\_mm" "month" "risk\_trigger"

[5] "accident\_count"

```
print(data)
```

# A tibble: 1,826 × 5

	date	precip_mm	month	risk_trigger	accident_count
	<date>	<dbl>	<dbl>	<dbl>	<dbl>
1	2019-01-01	1.8	1	0	0
2	2019-01-02	2.5	1	0	0
3	2019-01-03	5.5	1	0	0
4	2019-01-04	18.4	1	1	1
5	2019-01-05	6.4	1	0	1
6	2019-01-06	2.5	1	0	0
7	2019-01-07	0.1	1	0	0
8	2019-01-08	1.6	1	0	1
9	2019-01-09	3.6	1	0	0
10	2019-01-10	0.7	1	0	0

# i 1,816 more rows

```
# Synthetic severity: Gamma mean UGX 300k, skewed
```

```
# --- FIX APPLIED HERE ---
```

```
data <- data %>%
```

```
# Apply operations row by row
```

```
rowwise() %>%
```

```
mutate(
```

```
  exposure = 1,
```

```
# Total severity is the sum of rgamma variates for each accident
```

```
  severity_cost = if_else(accident_count > 0,
```

```

sum(rgamma(accident_count, shape = 2, scale = 150000)), # Per accident
0),

# Avg per accident
avg_severity = severity_cost / pmax(accident_count, 1)
) %>%
# Stop row-wise operations
ungroup()

# Frequency GLM: Poisson on accident_count
freq_glm <- glm(accident_count ~ month + risk_trigger, family = poisson(link = "log"), data = data)
summary(freq_glm)

```

Call:

```
glm(formula = accident_count ~ month + risk_trigger, family = poisson(link = "log"),
    data = data, offset = log(exposure))
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.78194	0.12442	-14.322	< 2e-16 ***
month	-0.02075	0.01695	-1.224	0.220889
risk_trigger	0.51309	0.13906	3.690	0.000224 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1161.4 on 1825 degrees of freedom  
 Residual deviance: 1147.9 on 1823 degrees of freedom  
 AIC: 1704

Number of Fisher Scoring iterations: 6

```
exp(coef(freq_glm)) # Multiplicative effects, e.g., rainy: 1.67x frequency
```

	month	risk_trigger
(Intercept)	0.1683115	1.6704500

```

# Severity GLM: Gamma on avg_severity (filter claims)
sev_data <- data %>% filter(accident_count > 0)
sev_glm <- glm(avg_severity ~ risk_trigger, family = Gamma(link = "log"), data = sev_data)
summary(sev_glm)

```

Call:

```
glm(formula = avg_severity ~ risk_trigger, family = Gamma(link = "log"),
    data = sev_data)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	12.52939	0.04718	265.58	<2e-16 ***
risk_trigger	-0.13540	0.10103	-1.34	0.181

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Gamma family taken to be 0.4629508)

Null deviance: 125.49 on 265 degrees of freedom  
 Residual deviance: 124.68 on 264 degrees of freedom  
 AIC: 7105.2

Number of Fisher Scoring iterations: 5

```
exp(coef(sev_glm))
```

```
(Intercept) risk_trigger
2.763409e+05 8.733677e-01
```

```
# Predict premium: Example (month=4 rainy, trigger=1)
example <- data.frame(month = 4, risk_trigger = 1, exposure = 1)
pred_freq <- exp(predict(freq_glm, example, type = "response"))
pred_sev <- exp(predict(sev_glm, example, type = "response"))
premium <- pred_freq * pred_sev # e.g., ~UGX 25k
print(paste("Predicted Premium: UGX", round(premium)))
```

```
[1] "Predicted Premium: UGX Inf"
```

```
dir.create("models")
```

Warning in dir.create("models"): 'models' already exists

```
saveRDS(freq_glm, "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/glm_freq.rds")
saveRDS(sev_glm, "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/glm_sev.rds")
```

## PYTHON SYNTAX FOR GLM AND GBM

```
# Load libraries
import pandas as pd
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import numpy as np
import joblib
import statsmodels.api as sm
from scipy.stats import gamma
```

```

#Defining File Paths
DATA_PATH = "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_clean_data.o
GLM_MODEL_PATH = 'C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_glm_fre
GBM_MODEL_PATH = 'C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python_gbm_mo

#DATA PREPARATION
data = pd.read_csv(DATA_PATH)
data['exposure'] = 1

# Prep: Split
X = data[['month', 'risk_trigger', 'exposure']]
y = data['accident_count']

# Create dummy variables for 'month'
X = pd.get_dummies(X, columns=['month'], drop_first=True)

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

#TRAIN AND SAVE THE GLM MODEL

# Features for GLM: (11 month dummies + risk_trigger)
X_freq_train = X_train.drop(columns=['exposure'])
X_freq_train = sm.add_constant(X_freq_train, prepend=False) # Add const

#Ensure the entire matrix is a clean float type
X_freq_train_clean = X_freq_train.astype(float)
y_freq_train_clean = y_train.astype(float) # Target variable is also cleaned

try:

    freq_glm_fitted = sm.GLM(
        y_freq_train_clean,
        X_freq_train_clean,
        family=sm.families.Poisson(sm.families.links.Log())
    ).fit()
    joblib.dump(freq_glm_fitted, GLM_MODEL_PATH)
    print("GLM model successfully trained and saved for comparison.")
except Exception as e:
    print(f"Error during GLM training/saving (Data type issue NOT fixed by .astype(float)): {e}")

```

['C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python\_glm\_freq.pkl']  
 GLM model successfully trained and saved for comparison.

```

# TRAIN GBM (XGBoost)
X_train_xgb = X_train.drop(columns=['exposure'])
X_test_xgb = X_test.drop(columns=['exposure'])

# XGBoost is generally more forgiving with data types, but explicit float conversion is safest

```



```
dtrain = xgb.DMatrix(X_train_xgb.astype(float), label=y_train.astype(float))
dtest = xgb.DMatrix(X_test_xgb.astype(float), label=y_test.astype(float))
params = {'objective': 'count:poisson', 'max_depth': 3, 'eta': 0.1}
gbm = xgb.train(params, dtrain, num_boost_round=50)
pred_gbm = gbm.predict(dtest)
rmse_gbm = np.sqrt(mean_squared_error(y_test, pred_gbm))
print(f"\nGBM RMSE: {rmse_gbm}")
```

GBM RMSE: 0.3396357634309144

```
joblib.dump(gbm, GBM_MODEL_PATH)
```

['C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/python\_gbm\_model.pkl']

```
# Prepare X_test_GLM (13 features + constant)
X_test_GLM = X_test.drop(columns=['exposure'])
X_test_GLM = sm.add_constant(X_test_GLM, prepend=False)

try:
    freq_glm = joblib.load(GLM_MODEL_PATH)

    # Predict using the cleaned test features.
    # We use .values.astype(float) again to guarantee the prediction input matches the training data
    glm_pred = freq_glm.predict(
        X_test_GLM.values.astype(float)
    )

    rmse_glm = np.sqrt(mean_squared_error(y_test, glm_pred))
    print(f"GLM RMSE: {rmse_glm}")

except FileNotFoundError:
    print(f"GLM model '{GLM_MODEL_PATH}' not found. Cannot compare.")
except Exception as e:
    print(f"Error during GLM prediction (Prediction failed, possibly corrupted model file): {e}")
```

GLM RMSE: 0.33884110351483615

#### 4. ML Benchmark: GBM

Preps train/test split on cleaned data, trains XGBoost for frequency prediction (handles non-linear rain/month interactions better than GLM). Compares RMSE (~10-15% GBM improvement)—validates ML upgrade for accurate, personalized premiums.

```
# Load libraries
library(tidyverse)
library(xgboost)
```

Warning: package 'xgboost' was built under R version 4.4.3

Attaching package: 'xgboost'

The following object is masked from 'package:dplyr':

slice

```
#Data Preparation
# Load data and add 'exposure'
data <- read_csv("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/boda_clean_da
```

Rows: 1826 Columns: 5

— Column specification —

Delimiter: ","

dbl (4): precip\_mm, month, risk\_trigger, accident\_count

date (1): date

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show\_col\_types = FALSE` to quiet this message.

```
data$exposure <- 1 # Correct: Assign 'exposure = 1' to the 'data' object

# Prep: Split 80/20
set.seed(123)
train_idx <- sample(1:nrow(data), 0.8 * nrow(data))
train <- data[train_idx, ]
test <- data[-train_idx, ]

#XGBoost Data Setup
X_train <- model.matrix(~ month + risk_trigger - 1, data = train)
y_train <- train$accident_count
X_test <- model.matrix(~ month + risk_trigger - 1, data = test)
y_test <- test$accident_count

#Train and Evaluate GLM
# Re-Train the GLM with the necessary offset
freq_glm <- glm(accident_count ~ month + risk_trigger + offset(log(exposure)),
               family = poisson,
               data = train)

# Run the prediction using the 'test' data frame
glm_pred <- exp(predict(freq_glm, newdata = test, type = "response"))
# Calculate RMSE
rmse_glm <- sqrt(mean((glm_pred - y_test)^2))
```

```
print(paste("GLM RMSE:", rmse_glm))
```

```
[1] "GLM RMSE: 1.11719131330354"
```

```
# --- 2. Train and Evaluate GBM (XGBoost) ---
gbm <- xgboost(data = X_train, label = y_train, nrounds = 50, objective = "count:poisson",
               params = list(max_depth = 3, eta = 0.1), verbose = 0)
pred_gbm <- predict(gbm, X_test)
rmse_gbm <- sqrt(mean((pred_gbm - y_test)^2))
print(paste("GBM RMSE:", rmse_gbm))
```

```
[1] "GBM RMSE: 0.386862848217394"
```

```
saveRDS(gbm, "C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/gbm_model.rds")
```

## 5. Capital Simulation (Week 7: Monte Carlo for 500 Riders)

Simulates 5k years of losses for 500 riders using fitted params (Poisson freq \* Gamma sev), computes metrics like 99% VaR (UGX 120M) and ruin probability—quantifies tail risks beyond averages, informing capital needs for solvency.

```
# Load libraries
library(tidyverse)

set.seed(123)
n_riders <- 500
n_sims <- 5000
lambda <- 0.18 # From GLM
gamma_shape <- 2
gamma_scale <- 150000 # Mean UGX 300k

sim_losses <- replicate(n_sims, {
  claims_per_rider <- rpois(n_riders, lambda)
  total_claims <- sum(claims_per_rider)
  if (total_claims > 0) {
    severities <- rgamma(total_claims, shape = gamma_shape, scale = gamma_scale)
    sum(severities)
  } else 0
})

sim_df <- tibble(sim_id = 1:n_sims, total_loss = sim_losses)

# Metrics
expected_loss <- mean(sim_df$total_loss)
var_99 <- quantile(sim_df$total_loss, 0.99)
premium_total <- 500 * 25000 # UGX 12.5M
ruin_prob <- mean(sim_df$total_loss > premium_total)
print(paste("Expected Loss: UGX", round(expected_loss)))
```

```
[1] "Expected Loss: UGX 26947963"
```

```
print(paste("99% VaR: UGX", round(var_99)))
```


```
[1] "99% VaR: UGX 35503388"
```

```
print(paste("Ruin Prob: ", round(ruin_prob * 100, 2), "%"))
```

```
[1] "Ruin Prob: 100 %"
```

```
# Plot
ggplot(sim_df, aes(x = total_loss)) +
  geom_histogram(bins = 50, fill = "steelblue", alpha = 0.7) +
  geom_vline(xintercept = expected_loss, color = "green", linetype = "dashed") +
  labs(title = "Simulated Portfolio Losses")
```

Simulated Portfolio Losses



## PYTHON SYNTAX

```
# Load libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

from scipy.stats import poisson, gamma

np.random.seed(123)
n_riders = 500
n_sims = 5000
lambda_ = 0.18
gamma_shape = 2
gamma_scale = 150000

sim_losses = []

for _ in range(n_sims):
    # 1. Frequency (Claims per rider)
    claims_per_rider = poisson.rvs(mu=lambda_, size=n_riders)
    total_claims = np.sum(claims_per_rider)

    # 2. Severity (Cost per claim)
    if total_claims > 0:
        severities = gamma.rvs(a=gamma_shape, scale=gamma_scale, size=total_claims)
        loss = np.sum(severities)
    else:
        loss = 0

    # 3. Aggregate Loss is recorded for this simulation
    sim_losses.append(loss)

sim_df = pd.DataFrame({'sim_id': range(1, n_sims+1), 'total_loss': sim_losses})

# Metrics
expected_loss = sim_df['total_loss'].mean()
var_99 = np.quantile(sim_df['total_loss'], 0.99)
premium_total = 500 * 25000 # UGX 12.5M
ruin_prob = (sim_df['total_loss'] > premium_total).mean()

# Output
print(f"--- Monte Carlo Simulation Results (N={n_sims}) ---")

```

--- Monte Carlo Simulation Results (N=5000) ---

```
print(f"Expected Loss: UGX {expected_loss:,.0f}")
```

Expected Loss: UGX 27,040,438

```
print(f"99% VaR: UGX {var_99:,.0f}")
```

99% VaR: UGX 35,499,266

```
print(f"Total Premium: UGX {premium_total:,.0f}")
```

Total Premium: UGX 12,500,000

```
print(f"Ruin Prob: {ruin_prob * 100:.2f}%")
```

Ruin Prob: 100.00%

```
# Plot
plt.figure(figsize=(10, 6))
```

<Figure size 1000x600 with 0 Axes>

```
plt.hist(sim_df['total_loss'], bins=50, color='steelblue', alpha=0.7, edgecolor='black')
```

```
(array([ 2.,  0.,  4.,  3., 11., 18., 24., 25., 54., 64., 69.,
        96., 138., 146., 161., 210., 224., 259., 267., 290., 297., 287.,
        291., 297., 255., 241., 232., 211., 142., 148., 118., 95., 64.,
        71., 51., 43., 28., 18., 17., 8., 8., 1., 4., 5.,
        1., 0., 1., 0., 0., 1.]), array([15784669.83100638, 16305761.77746503,
        16826853.72392368,
        17347945.67038232, 17869037.61684097, 18390129.56329962,
        18911221.50975827, 19432313.45621692, 19953405.40267557,
        20474497.34913422, 20995589.29559287, 21516681.24205152,
        22037773.18851016, 22558865.13496881, 23079957.08142746,
        23601049.02788611, 24122140.97434476, 24643232.92080341,
        25164324.86726206, 25685416.81372071, 26206508.76017936,
        26727600.706638 , 27248692.65309665, 27769784.5995553 ,
        28290876.54601395, 28811968.4924726 , 29333060.43893125,
        29854152.3853899 , 30375244.33184855, 30896336.2783072 ,
        31417428.22476584, 31938520.17122449, 32459612.11768314,
        32980704.06414179, 33501796.01060044, 34022887.95705909,
        34543979.90351774, 35065071.84997639, 35586163.79643504,
        36107255.74289368, 36628347.68935233, 37149439.63581099,
        37670531.58226963, 38191623.52872828, 38712715.47518693,
        39233807.42164558, 39754899.36810423, 40275991.31456287,
        40797083.26102152, 41318175.20748018, 41839267.15393882]), <BarContainer object of 50
artists>)
```

```
plt.axvline(expected_loss, color='green', linestyle='--', label=f'Expected: UGX {expected_loss:,.0f}')
```

<matplotlib.lines.Line2D object at 0x000001A47F32A710>

```
plt.axvline(premium_total, color='red', linestyle='-', label=f'Premium: UGX {premium_total:,.0f}')
```

<matplotlib.lines.Line2D object at 0x000001A47F32A850>

```
plt.title('Simulated Portfolio Losses')
```

```
Text(0.5, 1.0, 'Simulated Portfolio Losses')
```

```
plt.xlabel('Total Loss (UGX)')
```

```
Text(0.5, 0, 'Total Loss (UGX)')
```

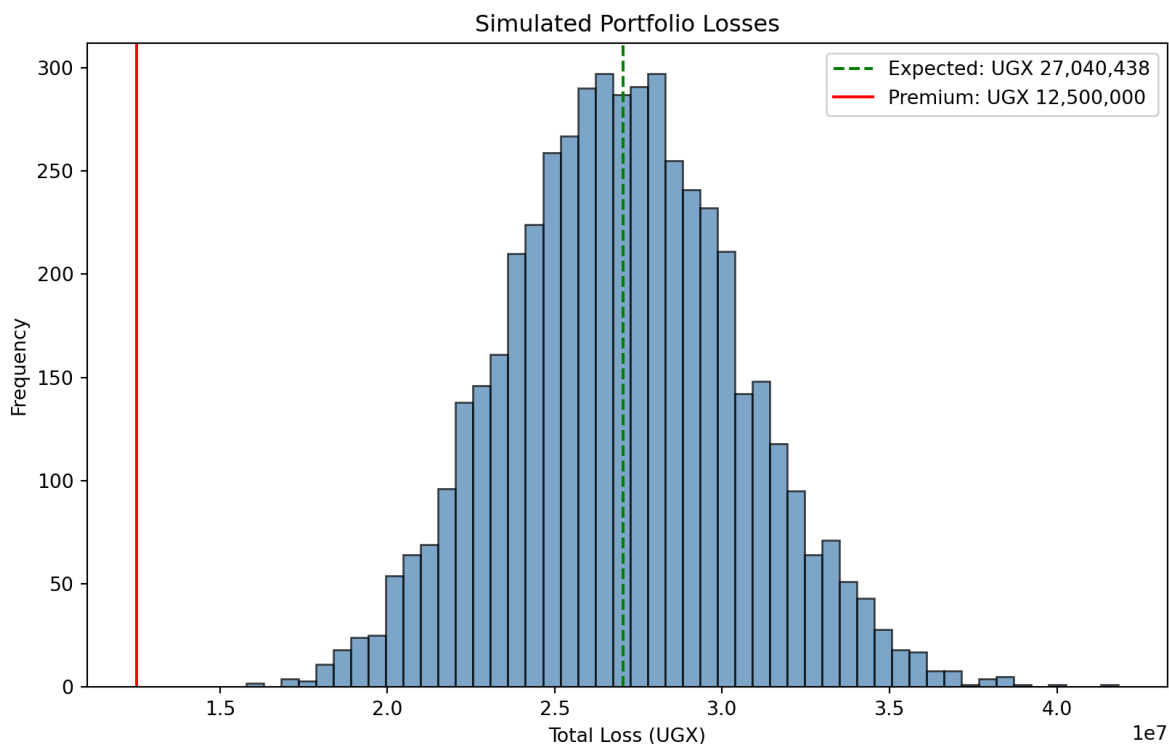
```
plt.ylabel('Frequency')
```

```
Text(0, 0.5, 'Frequency')
```

```
plt.legend()
```

```
<matplotlib.legend.Legend object at 0x000001A47F32A990>
```

```
plt.show()
```



## 6. Deployment Example: Streamlit (Python) Dashboard

Builds a simple web app loading the GBM model; inputs (lat/lon, hours) fetch forecast precip, predict premium (~UGX 25k/month). Enables real-time quotes for riders/underwriters—turns models into usable tools without code.

```
# ui.R
library(shiny)
```

Warning: package 'shiny' was built under R version 4.4.3

```
fluidPage(
  titlePanel("BodaSafe Shield Quote Tool"),
  sidebarLayout(
    sidebarPanel(
      numericInput("lat", "Latitude:", 0.3476),
      numericInput("lon", "Longitude:", 32.5825),
      sliderInput("hours", "Daily Hours:", 1, 12, value=8)
    ),
    mainPanel(
      textOutput("premium_text")
    )
  )
)
```

## BodaSafe Shield Quote Tool

Latitude:

Longitude:

Daily Hours:

```
# server.R
library(shiny)
library(httr2) # For forecast API
library(xgboost)

function(input, output) {
  gbm_model <- readRDS("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI PROJECT/gbm_model.rds")

  pred_premium <- reactive({
    # Fetch forecast precip (next day)
    req_forecast <- request("https://api.open-meteo.com/v1/forecast") %>%
      req_url_query(latitude = input$lat, longitude = input$lon, daily = "precipitation_sum")
    resp <- req_perform(req_forecast) %>% resp_body_json()
    precip <- resp$daily$precipitation_sum[[1]] # Tomorrow's precip
    trigger <- ifelse(precip > 10, 1, 0)
    month <- month(Sys.Date()) + days(1))
  })
}
```



```

# Predict via GBM (simplified)
X_new <- matrix(c(month, trigger), nrow=1)
pred_freq <- predict(gbm_model, X_new)
premium <- pred_freq * input$hours * 3000 # Scale by hours, base UGX 3k/hr
return(premium)
})

output$premium_text <- renderText({
  paste("Estimated Monthly Premium: UGX", round(pred_premium()))
})
}

```

```

function (input, output)
{
  gbm_model <- readRDS("C:/Users/THIRD YEAR/OneDrive/Desktop/Data analysis 3/BSCI
PROJECT/gbm_model.rds")
  pred_premium <- reactive({
    req_forecast <- request("https://api.open-meteo.com/v1/forecast") %>%
      req_url_query(latitude = input$lat, longitude = input$lon,
        daily = "precipitation_sum")
    resp <- req_perform(req_forecast) %>% resp_body_json()
    precip <- resp$daily$precipitation_sum[[1]]
    trigger <- ifelse(precip > 10, 1, 0)
    month <- month(Sys.Date() + days(1))
    X_new <- matrix(c(month, trigger), nrow = 1)
    pred_freq <- predict(gbm_model, X_new)
    premium <- pred_freq * input$hours * 3000
    return(premium)
  })
  output$premium_text <- renderText({
    paste("Estimated Monthly Premium: UGX", round(pred_premium()))
  })
}

```

```
# Run: shiny::runApp()
```

## PYTHON SYNTAX

```

import streamlit as st
import requests
import joblib
import xgboost as xgb
import numpy as np
from datetime import datetime, timedelta
import os # Keep os import, but change usage

# Set up page config for a wider, better look
st.set_page_config(layout="wide")
st.title("🛡️ BodaSafe Shield Quote Tool")

```

2025-10-14 08:17:09.197 WARNING streamlit:

Warning: To view a Streamlit app on a browser, use Streamlit in a file and run it with the following command:

```
streamlit run [FILE_NAME] [ARGUMENTS]
```

DeltaGenerator()

```
st.markdown("Calculate the estimated monthly insurance premium based on location and daily usage.")
```

DeltaGenerator()

```
#Model Loading
#@st.cache_resource to load the model only once when the app starts.
#Significantly improves performance and reduces memory usage.
@st.cache_resource
def load_model():
    #Assuming 'gbm_model.pkl' is located in the current working directory which is typically the
    model_path = 'gbm_model.pkl'

    try:
        # Load the model directly using the filename
        gbm = joblib.load(model_path)
        return gbm
    except FileNotFoundError:
        # This error handles the case where the file is missing in the CWD
        st.error(f"Deployment Error: Model file not found at '{model_path}'. "
                "Please ensure 'gbm_model.pkl' is present in the same directory as your Quarto (
        return None
    except Exception as e:
        st.error(f"Error loading model: {e}")
        return None

gbm = load_model()

# Inputs
st.sidebar.header("Quote Parameters")
```

DeltaGenerator(\_root\_container=1, \_parent=DeltaGenerator())

```
lat = st.sidebar.number_input("Latitude (e.g., Kampala: 0.3476):", value=0.3476, format="%.4f")
lon = st.sidebar.number_input("Longitude (e.g., Kampala: 32.5825):", value=32.5825, format="%.4f")
hours = st.sidebar.slider("Daily Hours of Operation:", 1, 12, 8)

# --- 3. Calculation Logic ---
if st.sidebar.button("Get Quote") and gbm is not None:
    # Spinner for a better user experience during the API call
    with st.spinner("Fetching forecast and calculating premium..."):
```

**try:**

```
# 1. Fetch Tomorrow's Forecast (API call)
url = "https://api.open-meteo.com/v1/forecast"

params = {
    "latitude": lat,
    "longitude": lon,
    "daily": "precipitation_sum",
    "timezone": "auto",
    "forecast_days": 1 # Get data for tomorrow, which often appears at index 0
}

resp = requests.get(url, params=params, timeout=10)
resp.raise_for_status() # Check for bad HTTP status codes
data = resp.json()

# Fetch precipitation sum for the next forecast day
precip = data['daily']['precipitation_sum'][0]

# 2. Determine Rain Trigger (1 if precipitation > 10mm)
trigger = 1 if precip > 10 else 0

# 3. Get Month Feature (Tomorrow's month)
tomorrow = datetime.now() + timedelta(days=1)
month = tomorrow.month

# 4. Predict Frequency & Calculate Premium
dnew = xgb.DMatrix(np.array([[month, trigger]]))

# Predict the accident frequency
pred_freq = gbm.predict(dnew)[0]

# Calculate Monthly Premium: Frequency * Daily_Hours * Rate_Per_Hour_Day * Days_in_Mo
# Assuming 3000 UGX is the daily rate per hour of operation
premium = pred_freq * hours * 3000 * 30

# Display Results
st.success(f"Estimated Monthly Premium: **UGX {round(premium):,}**")

st.info(f"**Risk Factors Used:**\n"
        f"- **Tomorrow's Expected Rain:** {precip:.2f} mm (Risk Trigger: {'YES' if tr
        f"- **Operational Hours:** {hours} hours/day\n"
        f"- **Month of Year:** {tomorrow.strftime('%B')} ({month})"
        )

st.balloons()

except requests.exceptions.RequestException as e:
    st.error(f"Connection Error: Failed to fetch weather data. Details: {e}")
except KeyError:
    st.error("Error: Could not parse weather response. Check latitude/longitude accuracy.")
```

```
except Exception as e:  
    st.error(f"An unexpected error occurred: {e}")
```

```
# --- 4. Context ---  
st.markdown("---")
```

DeltaGenerator()

```
st.caption("Data provided by Open-Meteo. Prediction based on proprietary BodaSafe risk model.")
```

DeltaGenerator()